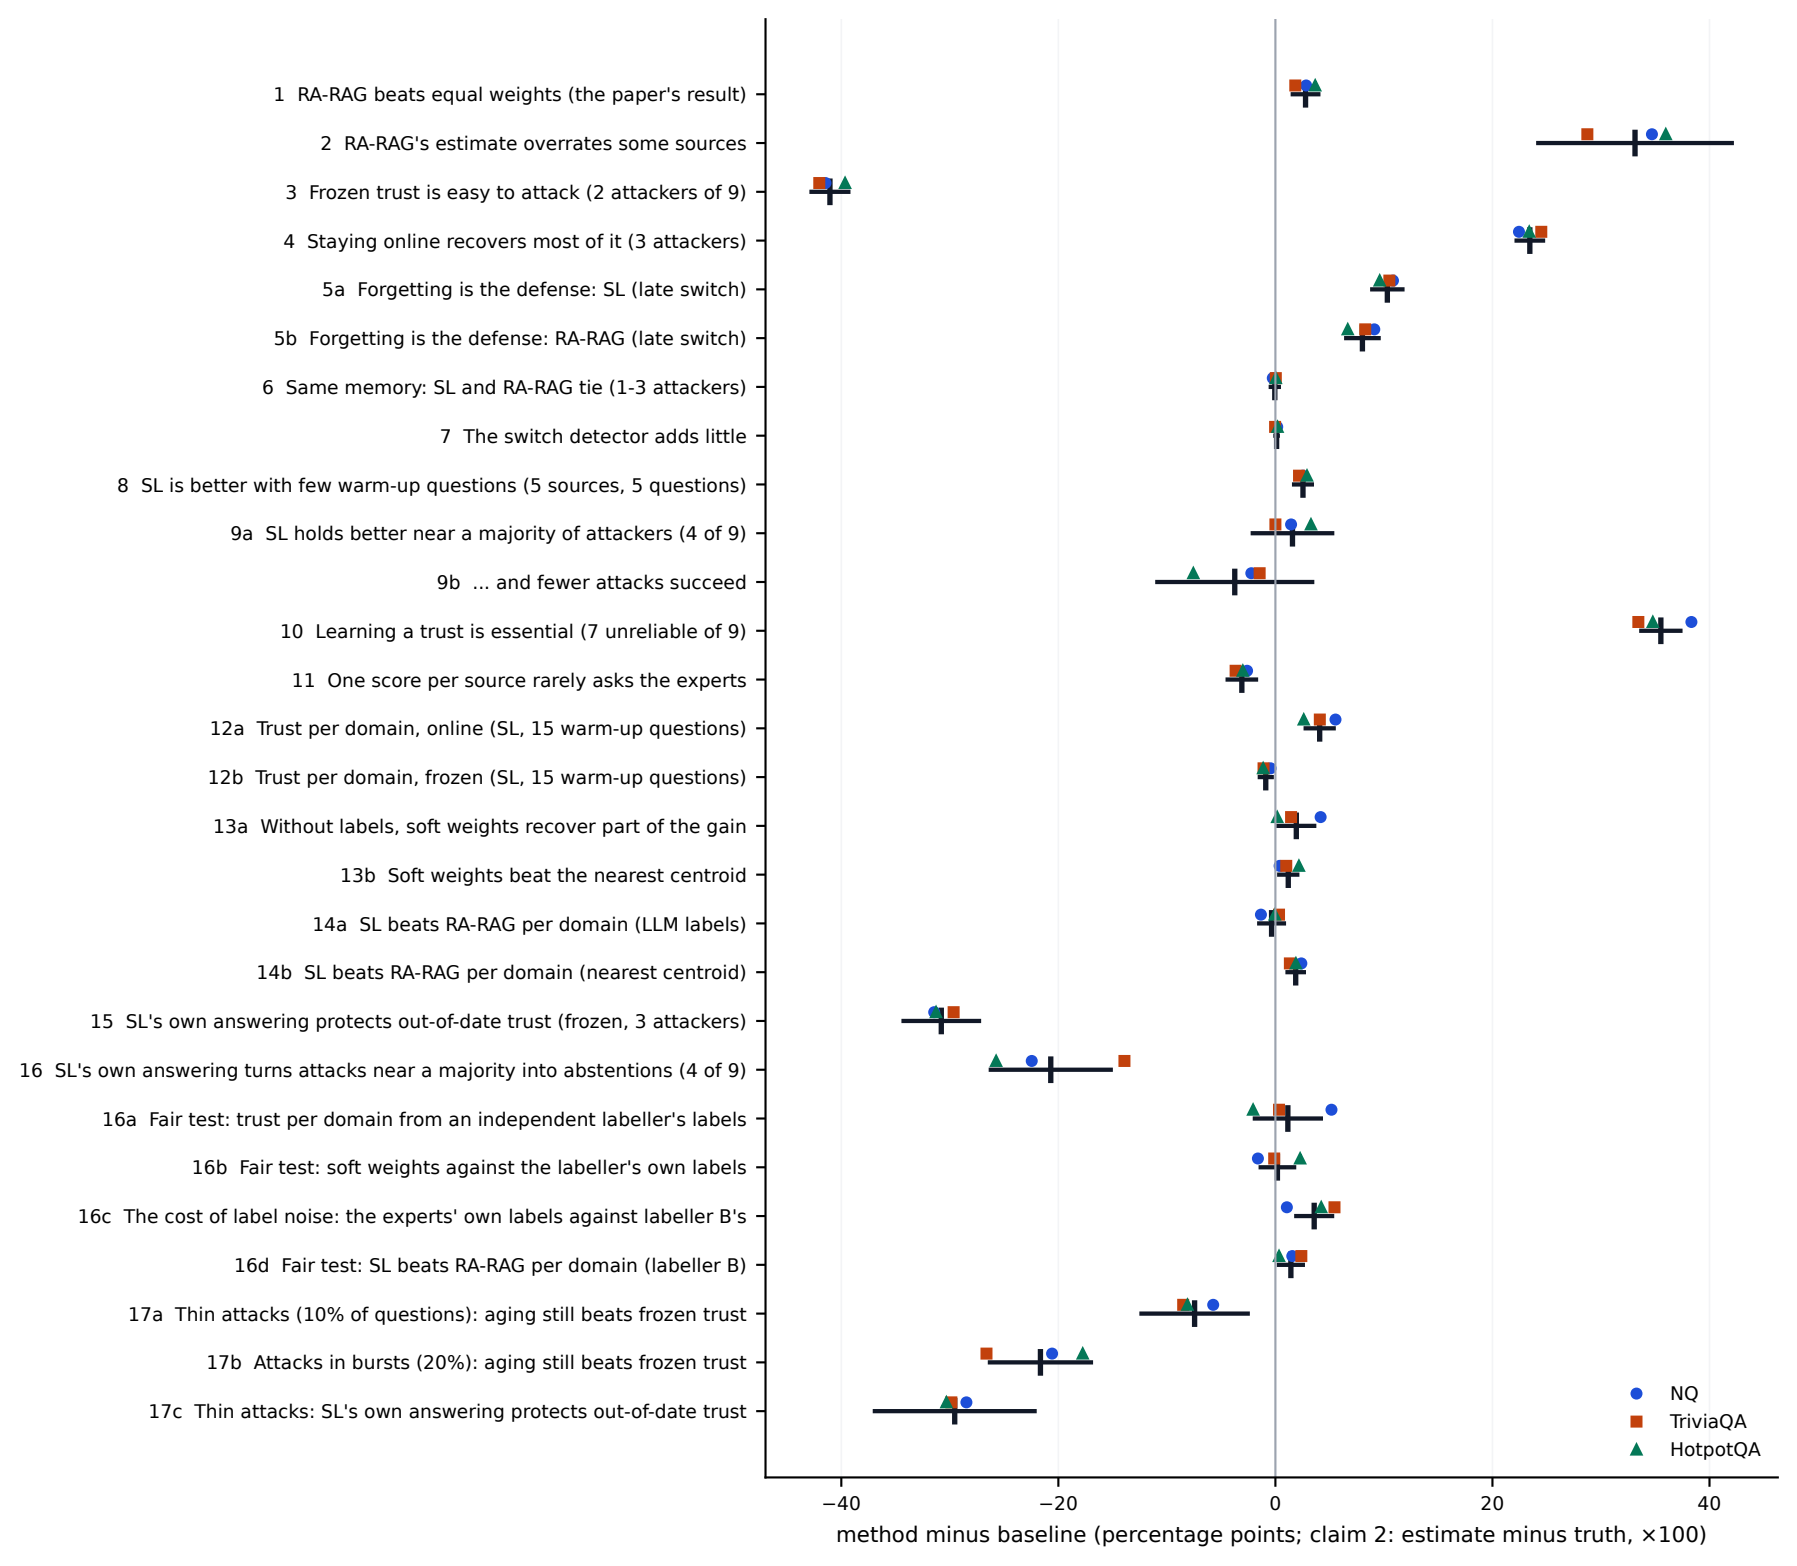
1 RA-RAG beats equal weights (the paper's result)
2 RA-RAG's estimate overrates some sources
3 Frozen trust is easy to attack (2 attackers of 9)
4 Staying online recovers most of it (3 attackers)
5a Forgetting is the defense: SL (late switch)
5b Forgetting is the defense: RA-RAG (late switch)
6 Same memory: SL and RA-RAG tie (1-3 attackers)
7 The switch detector adds little
8 SL is better with few warm-up questions (5 sources, 5 questions)
9a SL holds better near a majority of attackers (4 of 9)
9b ... and fewer attacks succeed
10 Learning a trust is essential (7 unreliable of 9)
11 One score per source rarely asks the experts
12a Trust per domain, online (SL, 15 warm-up questions)
12b Trust per domain, frozen (SL, 15 warm-up questions)
13a Without labels, soft weights recover part of the gain
13b Soft weights beat the nearest centroid
14a SL beats RA-RAG per domain (LLM labels)
14b SL beats RA-RAG per domain (nearest centroid)
15 SL's own answering protects out-of-date trust (frozen, 3 attackers)
16 SL's own answering turns attacks near a majority into abstentions (4 of 9)
16a Fair test: trust per domain from an independent labeller's labels
16b Fair test: soft weights against the labeller's own labels
16c The cost of label noise: the experts' own labels against labeller B's
16d Fair test: SL beats RA-RAG per domain (labeller B)
17a Thin attacks (10% of questions): aging still beats frozen trust
17b Attacks in bursts (20%): aging still beats frozen trust
17c Thin attacks: SL's own answering protects out-of-date trust
NQ
TriviaQA
HotpotQA
−40
−20
0
20
40
method minus baseline (percentage points; claim 2: estimate minus truth, ×100)